



# **AutoVision**

XLR8 PROJECT - ITSG 2025-2026

**PRESENTERS:** DAVID SZILAGYI, RAZVAN FILEA,  
ARMIN TOROK

# THE CHALLENGE: REAL-TIME INTELLIGENT PERCEPTION

## THE CONTEXT




REAL-TIME  
RESPONSIVENESS

SLOW =  
DANGEROUS

Simultaneous detection,  
tracking, recognition.

## THE PROBLEM



DEEP LEARNING  
(CNNs)      COMPUTATIONALLY  
EXPENSIVE



DELAY

SYNCHRONIZATION  
DELAYS & OVERHEAD

## THE TRADE-OFF



RAPID  
PROTOTYPING  
(Python)

DEPLOYMENT  
EFFICIENCY  
(C++/Rust)

# Initial ML Approach: Fine-Tuning on Real-World Data (BDD100K)




Figure 5.2: Training and validation loss evolution along with precision, recall, and mAP metrics across epochs. The late-epoch fluctuation corresponds to mosaic augmentation being disabled.

- **Methodology:**
  - Fine-tuned YOLO 11s (COCO-pretrained) on the **Berkeley DeepDrive 100K** dataset.
  - **Training:** 20 epochs, AdamW optimizer,  $640 \times 640$  resolution, focused on traffic lights and signs.
- **Theoretical Performance:**
  - Achieved stable convergence with validation **mAP@0.5** of **0.38**.
  - High precision for vehicles (**0.90**) and moderate precision for lights/signs in urban images. (**0.53, 0.56**)
- **The "Sim-to-Real" Failure:**
  - **Critical Domain Gap:** The model failed when deployed in the 1:10 scale in-office environment.
  - **False Positives:** Background objects (chairs, fire extinguishers) were frequently misclassified as traffic signals.



# Expectation meets Reality: The Sim-to-Real Gap



Figure 5.6: Qualitative detection results on the BDD100K test set using the fine-tuned YOLO11s model.





Figure 5.7: Detection performance in the real test environment.

# The Pivot to Sensor-Native Data

- Data Strategy:
  - Abandoned large-scale external datasets in favor of a **Handcrafted Dataset** captured directly from the vehicle's onboard sensors.
  - **Data Collection:** Manual (RC-mode) driving sessions to capture exact lighting and camera angles.
- Data Pipeline:
  - Applied custom augmentations: Brightness jitter, Gaussian noise, and rotation to simulate sensor noise and diversity .





# Model Specialization

- Architectural Change:
  - Replaced the monolithic model with **three specialized YOLO models** + a **geometric Lane Detector**
  - Why Specialization?
    - Direct Classification:** Detects states (e.g., Red/Green) immediately, eliminating the need for a secondary classification stage
    - No Compromises:** Avoids trade-offs in resolution and anchors by independently tuning for disparate objects
- Key Result:
  - Significantly mitigated background false positives and achieved robust detection of miniature props in real-time.







Figure 5.9: Qualitative detections in the scaled test environment after pretraining on the handcrafted dataset.

# The New Bottleneck: Concurrency & Latency



## The Consequence of Specialization:

We moved from 1 model to **4 concurrent perception branches**, forcing parallelism for performance



## Python's Limitation:

Standard Python multiprocessing struggles with **high-throughput video streams** .  
Leads to **high latency and memory overhead** .



## The Timing Requirement:

Fast branches (ex Lane Detection, ~6ms) must **update continuously** for steering control .  
Slow branches (ex Sign Recognition, ~28ms) must **not stall** the entire system .

# The Shared-Message Specification




Figure 4.1: Memory layout of SharedMessage: fields are aligned to enable atomic access to *Message Version*, while the *Mutex* and condition variables reside in shared memory.

Table 1. Write-Read Average Transfer Time (ms) 50.000 iterations

| Resolution | Size     | MP-Pipe | RS-IPC | SHM-LOCK | SHM-PTH |
|------------|----------|---------|--------|----------|---------|
| 256x256    | 576 KB   | 0.247   | 0.2    | 0.226    | 0.233   |
| 512x512    | 2.25 MB  | 0.908   | 0.728  | 0.796    | 0.82    |
| 640x480    | 2.64 MB  | 1.019   | 0.898  | 0.982    | 0.994   |
| 1280x720   | 7.91 MB  | 5.542   | 4.529  | 4.762    | 4.833   |
| 1920x1080  | 17.80 MB | 16.161  | 11.26  | 11.35    | 11.96   |


Table 2. Write-Read Total Transfer Time (s) 50.000 iterations

| Resolution | Size     | MP Pipe | RS-IPC  | SHM-LOCK | SHM-PTH |
|------------|----------|---------|---------|----------|---------|
| 256x256    | 576 KB   | 12.336  | 10.01   | 11.306   | 11.67   |
| 512x512    | 2.25 MB  | 45.411  | 36.408  | 39.789   | 40.986  |
| 640x480    | 2.64 MB  | 50.931  | 44.895  | 49.12    | 49.986  |
| 1280x720   | 7.91 MB  | 277.103 | 226.434 | 238.106  | 241.671 |
| 1920x1080  | 17.80 MB | 808.053 | 562.98  | 567.252  | 598.005 |



# Proposed Solution - Hybrid Architecture

- **Concept:** A "SharedMessage" infrastructure using Rust for low-level memory management
- **Design:**
  - **Python Layer:** Handles ML inference logging.
  - **Rust Layer:** Handles atomic synchronization, shared memory, and data flow and more
- **Benefit:** Enables high throughput without sacrificing the simplicity of Python development.



# The Core Algorithm - Partial-Blocking IPC

## The Communication Model.

- A Producer-Consumer strategy tailored for robotics.

## The Innovation: "Partial-Blocking" policy.

- Producer (Camera): Advances to the next frame once at least one consumer started processing it

## Why "Partial-Blocking"?

*Non-Blocking: Risks Desynchronization (input mismatch)*

Full-Blocking: Slowest model bottlenecks whole pipeline



Partial-Blocking: Ensures frame consistency while allowing fast branches to run at maximum speed

## The Outcome.


- Fast branches (lane detection) update frequently; slow branches (sign recognition) catch up asynchronously without stalling the system.

# Hybrid Python-Rust Architecture

- Design Philosophy:
  - **Python:** Retained for ML development, model loading, and visualization.
  - **Rust:** Handles low-level memory management, atomic synchronization, and data flow.
- Mechanism:
  - Modules communicate via a **Shared-Memory (IPC)** layer.
  - Ensures zero-copy data access where possible to maximize throughput.
- Benefit:
  - Allows models to be retrained/swapped in Python without modifying the compiled Rust infrastructure .




# DEMO






## Full-blocking Approach With standard IPC (~20FPS)



## Proposed Shared-Memory Partial-Blocking Architecture (~90FPS)



# QUANTITATIVE RESULTS: PERFORMANCE VALIDATION DIAGRAM

## NAIVE MULTIPROCESSING (Baseline)



20.1 FPS

THROUGHPUT (Workstation)




Slow, Bottlenecked



51.0 ms

PER-FRAME LATENCY



High Overhead, Blocking



TRANSFER TIME (1080p)

16.16 ms

Significant Delay

## HYBRID PYTHON-RUST (Ours)



86.9 FPS

THROUGHPUT (Workstation)




4.3x SPEEDUP, High-Flow



~14.1 ms

PER-FRAME LATENCY



REDUCED BY ~72%, Non-Blocking



TRANSFER TIME (1080p)



11.26 ms

30% FASTER IPC, Low Latency




✓ **CONCLUSION:** Real-Time Performance (35+ FPS) Achieved on Embedded Hardware (Jetson AGX Orin)

# New GUI





# SWOT ANALYSIS: HYBRID PERCEPTION ARCHITECTURE





# Conclusion & Future Work

- **Summary:**
  - We introduced a modular architecture combining Python's flexibility with Rust's safety and efficiency.
  - The **Partial-Blocking** strategy successfully balances throughput and consistency.
- **Impact:**
  - **Validated** real-time performance on embedded devices, providing a scalable **foundation** for robotics research
  - Companion Paper published in [2025 RRIA, Romanian Journal of Information Technology and Automatic Control](#)
- **Future Work:**
  - **Next-Gen Perception:** Future iterations will explore integrating **Vision-Language Models (VLMs)** for semantic reasoning and adapting the runtime for the complex control stacks of **Humanoid Robots**.

